
Toward Adaptive Snap-HyDE for Legal RAG: A Four-Benchmark Bottleneck Study

Hamza Iqbal

Washington University in St. Louis
hiqbal@wustl.edu

Hanson Li

Washington University in St. Louis
hanson.l@wustl.edu

Josh Li

Washington University in St. Louis
josh.l@wustl.edu

Abstract

Legal retrieval-augmented generation is usually summarized as one accuracy number, but legal QA failures are not one phenomenon. We study snap-conditioned HyDE across BarExamQA, HousingQA, LegalBench-SCALR, and CaseHOLD, using one full-corpus legal result and paired $N = 200$ diagnostics across four legal benchmarks to probe RAG bottlenecks. On full-corpus BarExamQA, `rag_snap_hyde` improves Gemma 4 26B from 78.08% to 81.17% (+3.09pp, McNemar $p = 0.0130$) and Gemma 4 E4B from 58.49% to 62.18% (+3.68pp, $p = 0.0129$). The remaining evidence is heterogeneous in a structured way: SCALR is candidate-depth sensitive then saturates at top-5, HousingQA improves with metadata-filtered retrieval, and CaseHOLD improves gold-option retrieval without yet converting that into a reliable answer lift. We interpret this heterogeneity as evidence for adaptive snap-HyDE: a controller that places HyDE-style reasoning on the active bottleneck. The resulting agenda is generalizable in form, even where current rows are $N = 200$: diagnose whether a legal task needs more candidates, better query generation, answer anchoring, metadata filtering, or stronger evidence-to-option conversion, then spend the reasoning budget there.

1 Introduction

Legal QA is a natural target for retrieval augmentation. Many questions depend on statutes, holdings, jurisdiction, procedural posture, and precise rule exceptions that may not be reliably stored in a model’s parameters. However, legal RAG is also a setting where more retrieved text can easily hurt. A passage can be topically related but legally incomplete; a holding can mention the same doctrine while being the wrong option; a state statute can be right in subject matter but wrong in jurisdiction. These errors are not just retrieval failures or generation failures. They are interactions between retrieval, reasoning, and answer selection.

Our initial project pursued a broader agentic legal RAG system: route the question, decompose the issue, retrieve per gap, judge sufficiency, and replan. That design matches legal research practice, but it is difficult to evaluate scientifically because any gain could come from retrieval, extra tokens, a second reasoning pass, or prompt luck. We therefore moved from full-system complexity to atomic retrieval interventions.

The central method family is snap-conditioned HyDE. A “snap” is a first-pass answer or reasoning trace. A HyDE-style hypothetical passage is then generated from that snap and used for retrieval. In the core sequential and two-call HyDE modes, the final answer sees the question and retrieved evidence; the snap shapes retrieval rather than being handed directly to final synthesis. Separate

snap-only and one-call ablations test what happens when visible snap reasoning is itself part of the final answer path. The intuition is legal: before asking the model to read external authority, ask it to state the issue and likely rule. Retrieval is then shaped by that legal frame rather than by the raw question alone.

The current legal-only evidence varies by task, but it varies in a structured way. BarExamQA shows a real full-corpus snap-HyDE lift, but top-1 versus top-5 retrieval is flat, so the lift should not be described as “more documents help.” LegalBench-SCALR shows a strong top-1 collapse and top-5 saturation, which suggests candidate-set size matters more than pseudo-document query formulation. HousingQA has a promising top-10 signal and, after fixing a metadata-casing bug, state-filtered retrieval reaches 62.5% at $k=10$, beating generic top-5 retrieval and directionally beating generic top-10. CaseHOLD now has repaired gold-option retrieval, and snap-HyDE raises gold retrieval from 16.0% to 47.0%, but answer accuracy only moves from 69.5% to 72.0% with $p = 0.4421$.

We exclude MuSiQue from the main report because it is not a legal benchmark. The final class report is framed around four legal tasks; non-legal multi-hop QA remains only an internal mechanism check. The forward hypothesis is that snap-HyDE should become adaptive: it should choose whether to spend its reasoning budget on query framing, query diversity, candidate expansion, metadata filtering, or answer-option conversion depending on what the problem appears to need. This is the generalization claim we can defend now: not that one prompt wins every benchmark, but that legal RAG errors can be typed into reusable bottlenecks with reusable interventions.

2 Related Work

Retrieval-augmented generation combines parametric generation with retrieved non-parametric evidence [Lewis et al., 2020]. HyDE generates a hypothetical document and embeds it for zero-shot retrieval [Gao et al., 2023], while Query2Doc expands queries with LLM-generated pseudo-documents [Wang et al., 2023]. FLARE retrieves during generation by using predicted future content as a query [Jiang et al., 2023]. Our method family inherits directly from this line: generated text is used as a retrieval handle, not as trusted evidence.

Adaptive RAG work asks when or how retrieval should be invoked. Self-RAG learns to retrieve and critique through reflection tokens [Asai et al., 2024]. Corrective RAG grades retrieved evidence and triggers correction when retrieval is unreliable [Yan et al., 2024]. Adaptive-RAG routes questions by complexity [Jeong et al., 2024]. Speculative RAG drafts multiple answers from retrieved subsets and verifies them with a larger model [Wang et al., 2025]. These systems motivate our future direction, but we do not claim to implement them here. Instead, we use their shared lesson: retrieval should be conditional on observed failure mode.

Legal NLP benchmarks increasingly show that legal RAG cannot be reduced to generic open-domain QA. LegalBench provides a broad suite of legal reasoning tasks, including SCALR holding selection [Guha et al., 2023]. CaseHOLD evaluates whether a model can select the holding referenced by a court-opinion context [Zheng et al., 2021]. Zheng et al. introduce BarExamQA and Housing Statute QA to evaluate realistic legal retrieval and answering workflows [Zheng et al., 2025]. LegalBench-RAG focuses on retrieving minimal relevant legal snippets [Pipitone and Alami, 2024], and Legal RAG Bench provides an end-to-end legal RAG benchmark with factorial retriever/generator decomposition [Butler and Butler, 2026]. RAGChecker similarly argues for fine-grained separation of retrieval and generation errors [Ru et al., 2024]. Our work follows this diagnostic turn: the same RAG method should be judged differently depending on whether the active failure is retrieval depth, candidate-set selection, metadata filtering, or evidence conversion.

We also considered hypothetical-representation work beyond HyDE. HyQE ranks contexts with hypothetical query embeddings [Zhou et al., 2024]. We did not implement a separate “HyRE” or hypothetical reasoning embedding index in this report. The implemented object is simpler: online snap-conditioned HyDE used as a legal retrieval probe.

3 Framework Design

3.1 Method ladder

The report uses an inheritance-style method ladder. Each method adds one capability where possible, so negative rows remain interpretable.

Method	What it isolates
llm_only	Parametric legal answer without retrieval.
snap_only_in_final	Whether a first-pass answer/reasoning trace helps without retrieval.
rag_simple	Standard retrieval plus answer generation.
rag_hyde	Question-conditioned hypothetical document retrieval.
rag_snap_hyde	Snap-conditioned HyDE retrieval plus final synthesis.
rag_snap_hyde_1call	Whether snap reasoning alone after ordinary retrieval is enough.
rag_snap_hyde_2call	Fixed-cost snap and HyDE passage in one first call, final synthesis in a second.
multi_hyde_diverse	Multiple generated retrieval passages for query diversity.
subagent_rag	Gap-routed report generation; currently used as a negative over-abstention control.

3.2 Snap-conditioned HyDE

Figure 1 shows the core method. The model first writes a snap answer or legal analysis. It then writes a hypothetical passage that states the likely controlling rule, fact, or holding. That passage is embedded and used for retrieval. The final model sees retrieved passages along with the original question; the snap is mediated through the retrieval query unless the run is an explicit snap-visible ablation.




Figure 1: Snap-conditioned HyDE. The generated passage shapes retrieval; the retrieved corpus passages remain the external evidence used by final synthesis.

3.3 Bottleneck diagnostics

We use three families of diagnostics. First, top- k sensitivity tests whether answer accuracy depends on retrieving more than the single best passage. This is a retrieval-policy stress test, not pure context truncation, because the cross-encoder reranks a first-stage candidate pool before selecting the final top- k . Second, snap-only and HyDE-only ablations ask whether the gain comes from answer anchoring, generated-query retrieval, or their combination. Third, gold-retrieval and disagreement audits ask whether better retrieval actually converts into better final answers.

The four legal benchmarks map naturally onto different bottleneck hypotheses: BarExamQA tests answer anchoring under strong legal priors, HousingQA tests state-specific legal retrieval, LegalBench-SCALR tests small candidate-set holding selection, and CaseHOLD tests whether a retrieved holding can be converted into the correct answer option. Keeping most diagnostic rows at $N = 200$ also makes the comparison deliberately parallel: the same slice size lets us ask which mechanism changes under a fixed evidence budget before scaling every dataset to full corpus.

4 Experimental Setup

Benchmarks. BarExamQA contains 1,195 bar-exam multiple-choice questions with fact patterns and legal passages. HousingQA contains yes/no housing-law questions over state statutes. LegalBench-SCALR contains 571 Supreme Court holding-selection questions. CaseHOLD contains




Figure 2: Grounded adaptive snap-HyDE schematic. The controller first observes cheap diagnostics, then chooses a route tied to the current benchmark signal: answer framing for BarExamQA, candidate expansion for SCALR, jurisdiction-aware metadata filtering for HousingQA, or option conversion for CaseHOLD.

court-opinion contexts and five candidate holdings. We deliberately exclude non-legal MuSiQue from the main benchmark set, although it remains useful for internal multi-hop stress testing.

Models. The full-corpus BarExamQA rows use Gemma 4 26B-A4B and Gemma 4 E4B served through the shared evaluation harness. LegalBench-SCALR and CaseHOLD diagnostic rows use Llama 70B. HousingQA uses Gemma 4 26B. For paired claims, we compare runs on the same question slice and provider. Claims are scoped to these model/provider settings rather than presented as model-independent benchmark rankings. Except for the full-corpus BarExamQA rows, we keep the newer legal diagnostics at $N = 200$ for parity, cost control, and fast paired auditing; the same harness can increase N once a route has enough signal to justify a larger run.

Metrics. We report accuracy for multiple-choice and yes/no datasets. Paired deltas use exact McNemar tests with b/c discordance counts. Retrieval diagnostics report gold retrieval only when the dataset’s gold mapping and corpus collection are known to be valid. The table rows and paired deltas in this draft were re-checked against detail logs in `reports/final_class_report/evidence_snapshot.md`.

5 Results and Analysis

5.1 BarExamQA: full-corpus snap-HyDE lift

BarExamQA is the strongest legal positive result. On Gemma 4 26B-A4B, `rag_snap_hyde` reaches 81.17%, compared with 78.08% for `rag_simple` (+3.09pp, b/c=124/87, $p = 0.0130$). On Gemma 4 E4B, the same method improves 58.49% to 62.18% (+3.68pp, b/c=172/128, $p = 0.0129$).

The mechanism is not pure retrieval depth. On a paired $N = 200$ diagnostic, `rag_simple` top-5 is 82.5% and top-1 is 83.0% ($p = 1.0$). The full-corpus lift is better read as snap-conditioned query anchoring plus evidence reconciliation, with model-size dependence: the 26B audit suggests the first-pass legal frame is already strong, while retrieval/synthesis appears more useful on the smaller E4B model. A legal MC model can have strong prior knowledge of the rule; retrieval then helps only if it is framed in a way that does not overwrite the right legal analysis with a partial snippet. Figure 3 shows this cross-size pattern.

Table 1: BarExamQA full-corpus Gemma 4 26B-A4B method matrix, $N = 1195$.

Method	Accuracy	Δ vs rag_simple
rag_simple	78.08%	–
golden_passage	78.66%	+0.58pp
rag_hyde	78.91%	+0.83pp
llm_only	79.75%	+1.67pp
snap_only_in_final	80.59%	+2.51pp
rag_snap_hyde	81.17%	+3.09pp




Figure 3: BarExamQA cross-size result. rag_snap_hyde improves over rag_simple on both Gemma 4 model sizes. The snap conditions retrieval; visible snap reasoning is tested by separate controls.

5.2 Four-benchmark legal bottleneck matrix

Table 2 is the compact result summary. It intentionally includes negative and directional rows because those rows identify where the next adaptive route should spend budget. The current evidence is therefore more promising as a bottleneck map than as a single leaderboard: it shows which legal bottleneck each dataset is probing.

Figure 4 summarizes the key diagnostic split: the same top- k or HyDE intervention has different meaning depending on whether accuracy and gold retrieval move together.

The most important negative result is CaseHOLD. After repairing the holdings corpus and gold mapping, top-1 rag_simple reaches 64.5% with 9.0% gold retrieval, while top-5 reaches 69.5% with 16.0% gold retrieval (+5.0pp, $p = 0.0525$). Top-10 retrieves more gold options (19.0%) but drops to 68.0%. HyDE retrieves the gold option in 52.0% of rows and two-call snap-HyDE in 47.0%, but both answer at 72.0%; neither is a significant lift over top-5 rag_simple. If retrieval recall were sufficient, answer accuracy should improve sharply. It does not. This suggests an evidence-to-option bottleneck: retrieved holdings can still be paraphrased distractors, overly generic, or hard to map back to the exact candidate label.

SCALR gives a different holding-selection story. Top-1 is poor, top-5 fixes most of the loss, top-10 adds no answer accuracy beyond top-5, and a validated rag_hyde follow-up is also flat at 77.5% despite retrieving the gold holding in 66.0% of rows. A one-call snap-final variant drops to 70.5% ($p = 0.00235$ versus top-5), so simply adding a snap rationale to the final answer is not a free improvement on this task. Diverse HyDE also stays below top-5 at 75.5%. This is not a generic “legal holdings ignore retrieval” result. It says SCALR needs a small candidate set, while CaseHOLD needs better evidence conversion.

Table 2: Legal-only benchmark matrix. MuSiQue is excluded because it is not legal-domain data.

Benchmark	Active bottleneck	Best current evidence	Caveat
BarExamQA / Gemma 4	Snap-conditioned query anchoring under strong legal priors	Full $N = 1195$: snap-HyDE is +3.09pp at 26B and +3.68pp at E4B; top-5 82.5% vs top-1 83.0% on $N = 200$	Not a pure retrieval-depth win.
HousingQA / Gemma 4	Depth, answer bias, and jurisdiction filtering	Top-1 50.5% to top-10 58.0% (+7.5pp, $p = 0.0722$); state filter reaches 61.5% at $k=5$ and 62.5% at $k=10$	State filter beats generic top-5; $k=10$ adds little over $k=5$.
LegalBench-SCALR / Llama 70B	Candidate-set size	Top-1 59.5%, top-5 77.0%, top-10 77.0%; HyDE 77.5%; diverse HyDE 75.5%	Depth helps to top-5; generated queries do not yet add answer lift.
CaseHOLD Llama 70B	Evidence-to-option conversion	Top-1 64.5%, top-5 69.5%, top-10 68.0%; HyDE and two-call both 72.0%	Better retrieval recall is not yet a reliable answer lift.




Figure 4: Legal benchmarks expose different bottlenecks. BarExamQA is depth-flat, SCALR benefits sharply from top-5 candidates, HousingQA improves with state-filtered retrieval, and CaseHOLD shows that gold-option retrieval can improve faster than answer accuracy.

HousingQA is the strongest evidence that legal metadata is not optional. The first state-filter experiment failed because question states were display-case strings while Chroma metadata was lower-case; that failed row is rejected as empty retrieval. After repairing the casing bug, $k=5$ state-filtered retrieval lands at 61.5%, with zero empty retrieval rows and 81/200 gold statutes retrieved. The completed $k=10$ state-filter run reaches 62.5%, with zero empty retrieval rows and 98/200 gold statutes retrieved. This is +9.0pp over generic top-5 rag_simple ($b/c=34/16$, $p = 0.015$) and +4.5pp over generic top-10 ($p = 0.306$), while adding only +1.0pp over state-filtered $k=5$ ($p = 0.815$). The claim is therefore not that deeper state-filtered retrieval is optimized; it is that the previous “Housing is blocked” story has become a concrete metadata-filtering signal.

5.3 BarExam single gold passage is a noisy control

The BarExamQA golden-passage control appears paradoxical: providing one gold passage reaches only 78.66%, below `llm_only` at 79.75% and `snap_only_in_final` at 80.59%. The reason is legal evidence granularity. A labeled passage can be topically correct but not decisive for the fact pattern; it can also highlight a distractor doctrine. Snap-first reasoning gives the model an initial legal frame before evidence reconciliation, which partly protects against over-reading one snippet. This claim is specific to the BarExamQA single-passage control; SCALR behaves differently, with

high conditional accuracy when the gold holding is retrieved. Figure 5 visualizes the paired flips behind this control.




Figure 5: BarExamQA mechanism audit. A single gold passage is a noisy control rather than an oracle ceiling; snap-first reasoning can recover some cases where forced context distracts the model.

5.4 Cost and budget behavior

The ablations also expose when extra calls and tokens are worth spending. Figure 6 plots paired $N = 200$ legal slices with marker size proportional to total tokens per question. BarExamQA is the cleanest case for snap-style reasoning budget: the two-call version reaches 85.5%, nearly matching the sequential snap-HyDE row at 86.0%, while using fewer calls and fewer tokens per question. HousingQA behaves differently: generic top-10 retrieval is one call but expensive in total tokens (4.6k tokens per question), while two-call snap-HyDE is slightly lower accuracy at a smaller 3.2k tokens per question. State-filtered retrieval is now the stronger Housing candidate because it uses jurisdiction metadata directly rather than spending generic depth; k=10 state filtering is 62.5%, only +1.0pp above k=5. The generic top-10 row also reduces a top-1 tendency to over-predict “Yes” on a dataset whose gold labels are mostly “No,” so the generic depth gain should not be over-described as jurisdiction repair. SCALR and CaseHOLD show the main caution: two-call reasoning roughly doubles calls and tokens, but SCALR does not improve and CaseHOLD improves only modestly. This is why we frame method choice as budgeted routing, not as always-on snap-HyDE.




Figure 6: Calls and tokens per question matter. Extra reasoning budget is useful on BarExamQA, ambiguous on HousingQA and CaseHOLD, and not currently justified on SCALR. Token values are computed from the paired detail logs.

5.5 Remaining validations

The newly landed validations should still be treated as diagnostic rows rather than final benchmark claims. HousingQA k=10 state filtering supports the metadata-filtering read: it improves

over generic top-10 directionally but does not materially improve over state-filtered $k=5$. SCALR follow-ups have also landed locally as flat or negative, so the next SCALR question is not more query generation but option grounding or reranking.

6 Discussion: From Method Search to Bottleneck Routing

The most useful interpretation of these results is not that snap-HyDE is the final method. It is that HyDE-style generation gives RAG a place to reason about the problem before retrieval, and each dataset reveals a different way that representation step should be aimed. Sometimes the right move is a more legally framed query; sometimes it is more query diversity, a larger candidate set, a metadata filter, or an answer-option mapper after retrieval. This matters because many RAG papers collapse several different phenomena into one comparison: retrieval versus no retrieval. In our logs, that comparison is too coarse. LegalBench-SCALR loses heavily at top-1 and recovers at top-5, so a route that expands the candidate set is plausible. BarExamQA is top- k flat but improves with snap-conditioned HyDE, so the route should preserve an answer frame rather than simply retrieve more passages. HousingQA likely requires metadata-constrained retrieval checks, because the question state is part of the legal problem and a normalized state filter now beats generic top-5 retrieval. CaseHOLD improves gold retrieval under two-call snap-HyDE but not answer accuracy, so the route should focus on mapping retrieved holdings back to answer options.

This gives a feasible version of the “agentic” idea without making the agent itself the unit of evaluation. The method we are working toward is an adaptive snap-HyDE controller: keep the snap/HyDE reasoning step, but decide how it should sit on top of the problem’s bottleneck. Instead of asking a planner to invent an open-ended research workflow, we can ask a controller to choose among a small set of evidence-budgeted interventions. The controller would observe cheap features before or during retrieval: question type, answer format, available metadata, top-score margin, retrieved-document agreement, candidate-option similarity, and whether the snap answer agrees with retrieval. It would then choose how to use the HyDE/reasoning budget: ordinary RAG when retrieval is already adequate, top- k expansion when candidate recall is binding, snap-conditioned HyDE when query framing is binding, diverse HyDE passages when one generated representation is too narrow, state-filtered retrieval when metadata is legally decisive, or option-aware reranking when retrieved evidence is not mapping back to choices. The result is still agentic in the sense that method adaptation happens on the fly, but it is experimentally cleaner because every route corresponds to a measured ablation cell.

Observed signal	Likely bottleneck	Route to test next
Top-1 fails, top-5 recovers, top-10 flat	Candidate-set size rather than answer reasoning	Use top-5 as default; spend budget on reranking or option grounding instead of deeper context.
Snap-only beats or matches retrieval	Evidence is distracting or under-specified relative to legal priors	Keep a first-pass legal frame in final synthesis and audit whether evidence overwrites correct priors.
Generated retrieval improves gold-hit but not answer accuracy	Evidence-to-answer conversion bottleneck	Add answer-option-aware synthesis, contradiction checks, or verifier scoring rather than more retrieval.
Question contains jurisdiction metadata	Corpus filtering and metadata normalization bottleneck	Filter before retrieval, validate non-empty evidence, and compare against generic top- k retrieval.
Retriever scores are low or retrieved passages disagree	Query formulation bottleneck	Try snap-HyDE or diverse HyDE passages, then measure whether gold retrieval and answer accuracy move together.

This framing also explains why the current negative results are useful. A negative CaseHOLD answer delta after a positive gold-hit delta is not just a failed method row; it identifies a downstream conversion problem. A flat BarExamQA top- k sweep does not make retrieval irrelevant; it says that retrieval depth is not the active lever for that slice. The HousingQA state-filter repair shows why metadata gates belong in the scientific evaluation path: a casing mismatch produced plausible-looking parametric answers with zero retrieved evidence, while the repaired filter changed

the method result. These are the hooks for a stronger final project: tune snap-HyDE quality where query formulation is binding, tune query number or diversity where one hypothetical passage is too narrow, tune metadata filters where jurisdiction is binding, and tune verifier or option-mapping logic where evidence conversion is binding. The stubborn cases are therefore not outside the framing; they define which HyDE-adjacent component still needs to be designed. The current evidence does not yet instantiate the full adaptive controller, but it shows why that controller is the right next object to build: different legal tasks need different forms of retrieval-time reasoning, and the harness can measure which form is helping.

Evidence-backed routes for adaptive snap-HyDE




Figure 7: Route map implied by current evidence. Green cells are supported by landed rows; amber cells are promising or next-test routes. The prospect for generalization comes from reusing the same bottleneck tests and routing choices across legal tasks, not from claiming one fixed prompt wins everywhere.

The class-report claim is therefore deliberately scoped but optimistic. We built and audited a legal RAG harness that can run comparable methods across several legal benchmarks; we found one full-corpus legal snap-HyDE lift on BarExamQA; we found several $N = 200$ diagnostic signals that expose reusable bottlenecks; and we turned those signals into a routing agenda. That is a better foundation for future work than another broad sweep over loosely comparable prompts, because it says what should change when a method fails.

7 Limitations and Future Work

The report is intentionally conservative about statistical scope. Only BarExamQA has a full-corpus legal result; HousingQA, SCALR, and CaseHOLD are paired $N = 200$ diagnostic slices for mech-

anism discovery rather than final benchmark claims. This was also a parity choice: the same sample size made it possible to compare mechanisms quickly under comparable API and audit budgets. The limitation is scale, not wiring. HousingQA now has repaired k=5 and k=10 state-filter results, but they should still be treated as diagnostics over one deterministic slice until we rerun larger samples.

The method set is incomplete relative to the literature. We define Self-RAG, CRAG, and Speculative RAG, but do not reimplement them or log true draft/verifier metrics. The next experiments should stay legal-only: finish test SCALR option grounding or reranking, test CaseHOLD option-aware synthesis after the gold-retrieval lift, and wire LegalBench-RAG or Legal RAG Bench rather than adding another generic multi-hop dataset. The adaptive snap-HyDE controller should be evaluated by scaling the strongest routes first: run larger HousingQA state-filter slices, add SCALR option reranking after top-5 retrieval, and test CaseHOLD option-aware conversion on the repaired holdings collection.

8 Conclusion

The legal-only evidence supports a bottleneck-typed view of RAG. Snap-HyDE helps BarExamQA at full corpus, but BarExamQA is not retrieval-depth limited. SCALR is candidate-depth limited, HousingQA shows a real metadata-filtering signal, and CaseHOLD shows that better generated-query retrieval does not automatically convert to better answers. This is the project signal: adaptive legal RAG should not ask whether HyDE is globally good or bad; it should ask where the task’s bottleneck is and apply the snap/HyDE reasoning budget there. That gives us a concrete and testable path from the current report to a more general method: diagnose the bottleneck, select the route, and scale the paired evaluation once the route has signal.

References

- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=hSyW5go0v8>.
- Abdur-Rahman Butler and Umar Butler. Legal RAG bench: An end-to-end benchmark for legal RAG. *arXiv preprint arXiv:2603.01710*, 2026. URL <https://arxiv.org/abs/2603.01710>.
- Luyu Gao, Xueguang Ma, Jimmy Lin, and Jamie Callan. Precise zero-shot dense retrieval without relevance labels. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, 2023. doi: 10.18653/v1/2023.acl-long.99. URL <https://aclanthology.org/2023.acl-long.99/>.
- Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Re, Adam Chilton, Alex Chohlas-Wood, Austin Peters, Brandon Waldon, Daniel N. Rockmore, Diego Zambrano, Sam Talisman, Crystal Ho, Roseanna Surani, Frank Fagan, Mona Sarfaty, Robert Dickinson, Niva Porat, Markus Hegland, Sijing Wu, et al. LegalBench: A collaboratively built benchmark for measuring legal reasoning in large language models. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2023. URL <https://arxiv.org/abs/2308.11462>.
- Soyeong Jeong, Jinheon Baek, Sukmin Cho, Sung Ju Hwang, and Jong Park. Adaptive-RAG: Learning to adapt retrieval-augmented large language models through question complexity. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 7036–7050, 2024. doi: 10.18653/v1/2024.naacl-long.389. URL <https://aclanthology.org/2024.naacl-long.389/>.
- Zhengbao Jiang, Frank F. Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu, Yiming Yang, Jamie Callan, and Graham Neubig. Active retrieval augmented generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023. doi: 10.18653/v1/2023.emnlp-main.495. URL <https://aclanthology.org/2023.emnlp-main.495/>.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen-tau Yih, Tim Rocktaschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*, 2020. URL <https://papers.nips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html>.
- Nicholas Pipitone and Ghita Houir Alami. LegalBench-RAG: A benchmark for retrieval-augmented generation in the legal domain. *arXiv preprint arXiv:2408.10343*, 2024. URL <https://arxiv.org/abs/2408.10343>.
- Dongyu Ru, Lin Qiu, Xiangkun Hu, et al. RAGChecker: A fine-grained framework for diagnosing retrieval-augmented generation. *arXiv preprint arXiv:2408.08067*, 2024. URL <https://arxiv.org/abs/2408.08067>.
- Liang Wang, Nan Yang, and Furu Wei. Query2doc: Query expansion with large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9414–9423, 2023. doi: 10.18653/v1/2023.emnlp-main.585. URL <https://aclanthology.org/2023.emnlp-main.585/>.
- Zilong Wang, Zifeng Wang, Long Le, Huaixiu Steven Zheng, Swaroop Mishra, Vincent Perot, Yuwei Zhang, Anush Mattapalli, Ankur Taly, Jingbo Shang, Chen-Yu Lee, and Tomas Pfister. Speculative RAG: Enhancing retrieval augmented generation through drafting. In *International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=xgQfWbV6Ey>.
- Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. Corrective retrieval augmented generation. *arXiv preprint arXiv:2401.15884*, 2024. URL <https://arxiv.org/abs/2401.15884>.
- Lucia Zheng, Neel Guha, Brandon R. Anderson, Peter Henderson, and Daniel E. Ho. When does pre-training help? assessing self-supervised learning for law and the CaseHOLD dataset of 53,000+ legal holdings. In *Proceedings of the Eighteenth International Conference on Artificial Intelligence and Law*, 2021. URL <https://arxiv.org/abs/2104.08671>.

Lucia Zheng, Neel Guha, Javokhir Arifov, Sarah Zhang, Michal Skreta, Christopher D. Manning, Peter Henderson, and Daniel E. Ho. A reasoning-focused legal retrieval benchmark. In *Proceedings of the ACM Symposium on Computer Science and Law*, 2025. URL <https://arxiv.org/abs/2505.03970>.

Weichao Zhou, Jiaxin Zhang, Hilaf Hasson, Anu Singh, and Wenchao Li. HyQE: Ranking contexts with hypothetical query embeddings. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024. URL <https://aclanthology.org/2024.findings-emnlp.761/>.